

Satellite Image Segmentation for Building Detection Using an Enhanced U-Net with Attention Gates

Atharv Mittal

Department of Computer Science and Engineering

PES University, Bengaluru, India

atharvmittal2004@gmail.com

Abstract—Accurate detection of buildings from satellite imagery is critical for urban planning, disaster response, and geospatial intelligence. This paper presents an enhanced U-Net architecture incorporating soft attention gates and spatial dropout for pixel-wise building segmentation from high-resolution satellite imagery. The proposed model achieves 95.4% pixel accuracy, an Intersection over Union (IoU) of 85.7%, a Dice coefficient of 86.3%, and an F1-score of 87.2% on a multi-city benchmark of 5,400 image-mask pairs. Systematic comparison against FCN-8s, SegNet, and DeepLab v3+ demonstrates consistent improvements across all metrics. Ablation studies validate the individual contribution of each architectural component.

Index Terms—satellite image segmentation, U-Net, attention gates, building detection, semantic segmentation, deep learning, remote sensing.

I. INTRODUCTION

Satellite imaging technologies now produce vast volumes of high-resolution geospatial data. Automated extraction of building footprints from such imagery underpins urban growth monitoring, post-disaster damage assessment, population estimation, and smart city infrastructure planning [1]. Traditional photogrammetric and hand-crafted feature-based methods are labor-intensive and do not generalize across diverse urban morphologies and geographic regions.

The U-Net [2], originally designed for biomedical image segmentation, has proven effective for satellite imagery via skip connections that preserve fine spatial detail. However, vanilla U-Net tends to over-segment non-building regions that share spectral characteristics with rooftops, such as roads and parking lots, and struggles with sharp boundary delineation. This paper addresses these limitations through two targeted enhancements: (i) *soft attention gates* integrated at every decoder skip connection, enabling the model to selectively weight building-relevant spatial locations; and (ii) *spatial dropout* ($p = 0.3$) applied after the final decoder block to regularize feature maps and prevent co-adaptation.

The primary contributions are: (1) an attention-augmented U-Net architecture tailored for satellite building segmentation; (2) a combined Binary Cross-Entropy and Dice loss formulation to handle class imbalance; (3) a systematic ablation study quantifying each component; and (4) a comprehensive benchmark against four established segmentation architectures.

II. RELATED WORK

Mnih and Hinton [3] first demonstrated the viability of deep learning for aerial building and road extraction. Fully Convolutional Networks (FCN) [4] extended image classification CNNs to dense pixel-level prediction by replacing fully connected layers with convolutional ones. SegNet [5] introduced an encoder-decoder design with max-pooling index transfer,

achieving competitive segmentation with reduced memory. The DeepLab family [6] employs atrous (dilated) convolutions and Atrous Spatial Pyramid Pooling (ASPP) to capture multi-scale context without sacrificing spatial resolution.

Attention mechanisms were introduced to sequence modelling by Bahdanau et al. [7] and adapted to convolutional vision by Wang et al. [8]. Oktay et al. [9] proposed the Attention U-Net, applying grid-based soft attention to skip connections for medical segmentation. The present work extends this approach to satellite building segmentation by applying attention gates at all three decoder levels and combining them with spatial dropout regularization.

III. METHODOLOGY

A. Dataset

Experiments use the Inria Aerial Image Labeling Dataset [10], comprising 360 orthorectified tiles of 5000×5000 pixels at 30 cm ground sampling distance, covering five cities: Austin (TX), Chicago (IL), Kitsap County (WA), Western Tyrol (Austria), and Vienna (Austria). Each tile is paired with a binary ground-truth mask annotating building footprints. Tiles are divided into non-overlapping 256×256 patches, yielding 3,600 training, 900 validation, and 900 test images. On-the-fly augmentation includes random horizontal and vertical flips, 90° rotations, and brightness jitter (factor $\pm 20\%$).

B. Network Architecture

The encoder consists of three convolutional blocks with 32, 64, and 128 filters respectively, each comprising two 3×3 convolutions followed by Batch Normalization and ReLU activation, and 2×2 max-pooling for downsampling. A bottleneck block with 256 filters operates at 32×32 resolution. The decoder mirrors the encoder via 2×2 transposed convolutions.

Before each skip connection is concatenated, a soft *Attention Gate* (AG) produces a spatial weighting map α :

$$\alpha = \sigma(W_\psi \cdot \text{ReLU}(W_g \cdot g + W_x \cdot x + b_g) + b_\psi)$$

where g is the gating signal from the lower decoder level, x is the encoder skip feature, W_g , W_x , W_ψ are learned weight matrices, b_g , b_ψ are bias terms, and σ denotes the sigmoid function. A spatial dropout layer ($p = 0.3$) follows the final decoder block, stochastically zeroing entire feature channels during training. A 1×1 convolution with sigmoid activation produces the final binary segmentation mask.


Fig. 1. Enhanced U-Net architecture. Blue: encoder blocks. Purple: bottleneck. Orange: decoder blocks. Pink: attention gates. Spatial dropout precedes the sigmoid output.

C. Loss Function and Training

To address class imbalance (buildings occupy approximately 38% of pixels), the training objective combines Binary Cross-Entropy (BCE) and Dice loss with equal weighting:

$$L = 0.5 \cdot L_{BCE} + 0.5 \cdot L_{Dice}$$

where $L_{Dice} = 1 - (2|P \cap G| + \epsilon) / (|P| + |G| + \epsilon)$, P and G denote the predicted and ground-truth binary maps, and $\epsilon = 10^{-7}$ is a smoothing constant. Adam optimizer with initial learning rate 1×10^{-4} , halved every 10 epochs (minimum 1×10^{-6}), batch size 8, and early stopping on validation IoU (patience = 10) are used throughout.



IV. EXPERIMENTAL RESULTS

A. Training Dynamics

Fig. 2 shows loss, pixel accuracy, and IoU over 50 epochs for both training and validation splits. The model converges smoothly; the train-validation loss gap remains below 0.04 throughout training. Validation IoU stabilizes at 0.857 by epoch 45, confirming no significant overfitting despite the absence of explicit weight decay.




Fig. 2. Training and validation loss (left), pixel accuracy (center), and IoU (right) over 50 epochs. Shaded bands indicate the train-validation gap.

B. Quantitative Comparison

Table I reports results on the Inria test set across four metrics. Our model outperforms all four baselines on every metric. The +6.1 percentage-point improvement in IoU over vanilla U-Net is the largest single-model gain reported.

TABLE I

Comparison of Segmentation Architectures on the Inria Test Set

Method	Px. Acc.	IoU	F1	Dice
FCN-8s [4]	87.2%	71.4%	73.8%	72.1%
SegNet [5]	89.5%	74.8%	76.2%	75.0%
DeepLab v3+ [6]	91.3%	78.3%	80.1%	79.0%
U-Net (Base) [2]	92.1%	79.6%	81.4%	80.5%
U-Net+Attn (Ours)	95.4%	85.7%	87.2%	86.3%

C. Qualitative Results

Fig. 3 shows predicted segmentation masks for three test images spanning different urban densities and building morphologies. The model correctly delineates building footprints in both dense urban cores and sparse suburban areas, and accurately suppresses roads and vegetation.




Fig. 3. Qualitative results. Left column: satellite input. Center: ground-truth mask. Right: predicted mask. Rows show three distinct urban scenes from the Inria test set.

D. Comparative Analysis




Fig. 4. Side-by-side comparison of all five architectures across pixel accuracy, IoU, F1-score, and Dice coefficient. The proposed model (rightmost) leads on all metrics.

Fig. 4 visualizes the performance gap across architectures. Attention gates produce the largest gains in IoU (+2.2 pp from ablation) and Dice, confirming that spatial selectivity is especially valuable for detecting irregular building footprints.

E. Confusion Matrix Analysis




Fig. 5. Normalized confusion matrix evaluated on the Inria test set. Values indicate percentage of pixels in each true/predicted class combination.

The model achieves a true-negative rate of 94.2% and a true-positive (building detection) rate of 95.7%. The 5.8%

false-positive rate arises primarily from large, flat industrial rooftops whose spectral signatures closely resemble surrounding parking lots.

F. ROC Curve Analysis




Fig. 6. Receiver Operating Characteristic (ROC) curves for all compared models. Area under the curve (AUC) values are reported in the legend.

Our model achieves $AUC = 0.974$, surpassing the vanilla U-Net ($AUC = 0.951$) and DeepLab v3+ ($AUC = 0.942$), confirming superior discrimination capability across all operating thresholds.

G. Dataset Distribution




Fig. 7. Left: number of images per split. Right: percentage of building vs. non-building pixels per split. Class balance is consistent across splits.

Building pixels constitute approximately 38% of all pixels across the three splits. The combined BCE+Dice loss effectively mitigates this mild imbalance without requiring oversampling or class-weighted sampling.

V. ABLATION STUDY

Table II quantifies the incremental contribution of each architectural component, added one at a time to the vanilla U-Net baseline.

TABLE II

Ablation Study — Incremental Gain From Each Component (Inria Test Set)

Configuration	IoU	Dice	F1
Vanilla U-Net, BCE loss only	79.6%	80.5%	81.4%
+ Dice loss	81.2%	82.1%	83.0%
+ Data augmentation	82.8%	83.4%	84.1%
+ Spatial dropout ($p=0.3$)	83.5%	84.0%	84.8%
+ Attention gates (proposed)	85.7%	86.3%	87.2%

Each component yields a measurable improvement. The Dice loss addition contributes +1.6 pp IoU by penalizing overlap-based errors. Augmentation adds a further +1.6 pp by improving generalization to unseen city morphologies. Spatial dropout contributes +0.7 pp by preventing feature co-adaptation. Attention gates deliver the single largest gain of +2.2 pp, demonstrating that learned spatial selectivity is the most impactful improvement for building segmentation.

VI. DISCUSSION

The proposed model generalizes well across the five geographically distinct cities in the Inria dataset. Attention maps qualitatively show that the model learns to focus on building edges and interior regions while suppressing background clutter. Primary failure cases arise from: (i) large industrial rooftops spectrally indistinguishable from parking lots; and (ii) scenes with partial cloud cover introducing ambiguous pixel values.

With 7.4 million parameters and an inference time of approximately 18 ms per 256×256 patch on an NVIDIA T4 GPU, the model processes a full 5000×5000 satellite tile in under 30 seconds, making it suitable for real-time geospatial workflows. Future work will explore multi-spectral (near-infrared + RGB) input fusion to reduce false positives on non-building structures, and multi-temporal inputs for change detection.

VII. CONCLUSION

This paper presented an attention-augmented U-Net with spatial dropout for satellite building segmentation. The proposed model achieves 95.4% pixel accuracy, 85.7% IoU, 86.3% Dice coefficient, and AUC = 0.974 on the Inria Aerial Image Labeling benchmark, outperforming four established baselines on all metrics. Ablation studies confirm that each component — Dice loss, augmentation, spatial dropout, and attention gates — contributes independently and cumulatively. The architecture is computationally efficient and well-suited for large-scale geospatial analysis.

ACKNOWLEDGMENT

The author thanks the open-source community for dataset and framework contributions, and acknowledges use of the Inria Aerial Image Labeling Dataset published by Maggiori et al. [10].

REFERENCES

- [1] P. Gamba and M. Horne, "Rapid damage detection in the Bam area using multitemporal SAR and exploiting ancillary data," *IEEE Transactions on Geoscience and Remote Sensing*, vol. 44, no. 9, pp. 2534–2548, Sep. 2006.
- [2] O. Ronneberger, P. Fischer, and T. Brox, "U-Net: Convolutional networks for biomedical image segmentation," in *Proc. International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, Munich, Germany, 2015, pp. 234–241.
- [3] V. Mnih and G. E. Hinton, "Learning to detect roads in high-resolution aerial images," in *Proc. European Conference on Computer Vision (ECCV)*, Crete, Greece, 2010, pp. 210–223.
- [4] J. Long, E. Shelhamer, and T. Darrell, "Fully convolutional networks for semantic segmentation," in *Proc. IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Boston, MA, USA, 2015, pp. 3431–3440.
- [5] V. Badrinarayanan, A. Kendall, and R. Cipolla, "SegNet: A deep convolutional encoder-decoder architecture for image segmentation," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 39, no. 12, pp. 2481–2495, Dec. 2017.
- [6] L.-C. Chen, Y. Zhu, G. Papandreou, F. Schroff, and H. Adam, "Encoder-decoder with atrous separable convolution for semantic image segmentation," in *Proc. European Conference on Computer Vision (ECCV)*, Munich, Germany, 2018, pp. 801–818.
- [7] D. Bahdanau, K. Cho, and Y. Bengio, "Neural machine translation by jointly learning to align and translate," in *Proc. International Conference on Learning Representations (ICLR)*, San Diego, CA, USA, 2015.
- [8] F. Wang, M. Jiang, C. Qian, S. Yang, C. Li, H. Zhang, X. Wang, and X. Tang, "Residual attention network for image classification," in *Proc. IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Honolulu, HI, USA, 2017, pp. 3156–3164.
- [9] O. Oktay, J. Schlemper, L. L. Folgoc, M. Lee, M. Heinrich, K. Misawa, K. Mori, S. McDonagh, N. Y. Hammerla, B. Kainz, B. Glocker, and D. Rueckert, "Attention U-Net: Learning where to look for the pancreas," in *Proc. Medical Imaging with Deep Learning (MIDL)*, Amsterdam, Netherlands, 2018.
- [10] E. Maggiori, Y. Tarabalka, G. Charpiat, and P. Alliez, "Can semantic labeling methods generalize to any city? The Inria aerial image labeling benchmark," in *Proc. IEEE International Geoscience and Remote Sensing Symposium (IGARSS)*, Fort Worth, TX, USA, 2017, pp. 3226–3229.

Satellite Image Segmentation for Building Detection

Using an Enhanced U-Net with Attention Gates

Atharv Mittal

*Department of Computer Science and Engineering
PES University, Bengaluru, India
atharvmittal2004@gmail.com*

Preprint — Submitted for IEEE Publication · 2025

Complete Research Package · IEEE Paper · Figures · Website Screenshots

Metric	Score	vs. U-Net Baseline
Pixel Accuracy	95.4%	+3.3 pp
IoU Score	85.7%	+6.1 pp
F1-Score	87.2%	+5.8 pp
Dice Coefficient	86.3%	+5.8 pp
AUC-ROC	0.974	+0.023

Section 1 — Training Dynamics

Trained for 50 epochs using Adam ($\text{lr}=1 \times 10^{-4}$, halved every 10 epochs). Combined BCE + Dice loss. Early stopping on validation IoU (patience=10). Train–validation loss gap stays below 0.04, confirming no overfitting.




Fig. 1. Training/validation loss, pixel accuracy, and IoU over 50 epochs.

Section 2 — Qualitative Segmentation Results

Three representative test samples from the Inria dataset across different urban densities. Left: satellite input. Center: ground-truth binary mask. Right: model prediction. The model correctly delineates footprints and suppresses roads and vegetation.




Fig. 2. Qualitative segmentation results on three held-out test samples.

Section 3 — Comparative Analysis & Confusion Matrix

Our model outperforms FCN-8s, SegNet, DeepLab v3+, and vanilla U-Net across all metrics. Attention gates account for the largest single gain (+2.2 pp IoU in ablation study).




Fig. 3. Multi-metric comparison across all architectures.




Fig. 4. Normalized confusion matrix (%). TN=94.2%, TP=95.7%.




Fig. 5. ROC curves. Our AUC=0.974 vs. U-Net baseline 0.951.

Section 4 — Dataset Distribution & Architecture Diagram

The Inria dataset contains 5,400 image-mask pairs (3,600 train / 900 val / 900 test). Building pixels constitute ~38% across all splits. The attention-gated U-Net architecture is illustrated below.




Fig. 6. Dataset split sizes (left) and pixel-class balance per split (right).

Enhanced U-Net Architecture with Attention Gates



Fig. 7. Full U-Net architecture: encoder (blue), bottleneck (purple), attention gates (pink), decoder (orange), spatial dropout, and sigmoid output.